



DATABRICKS

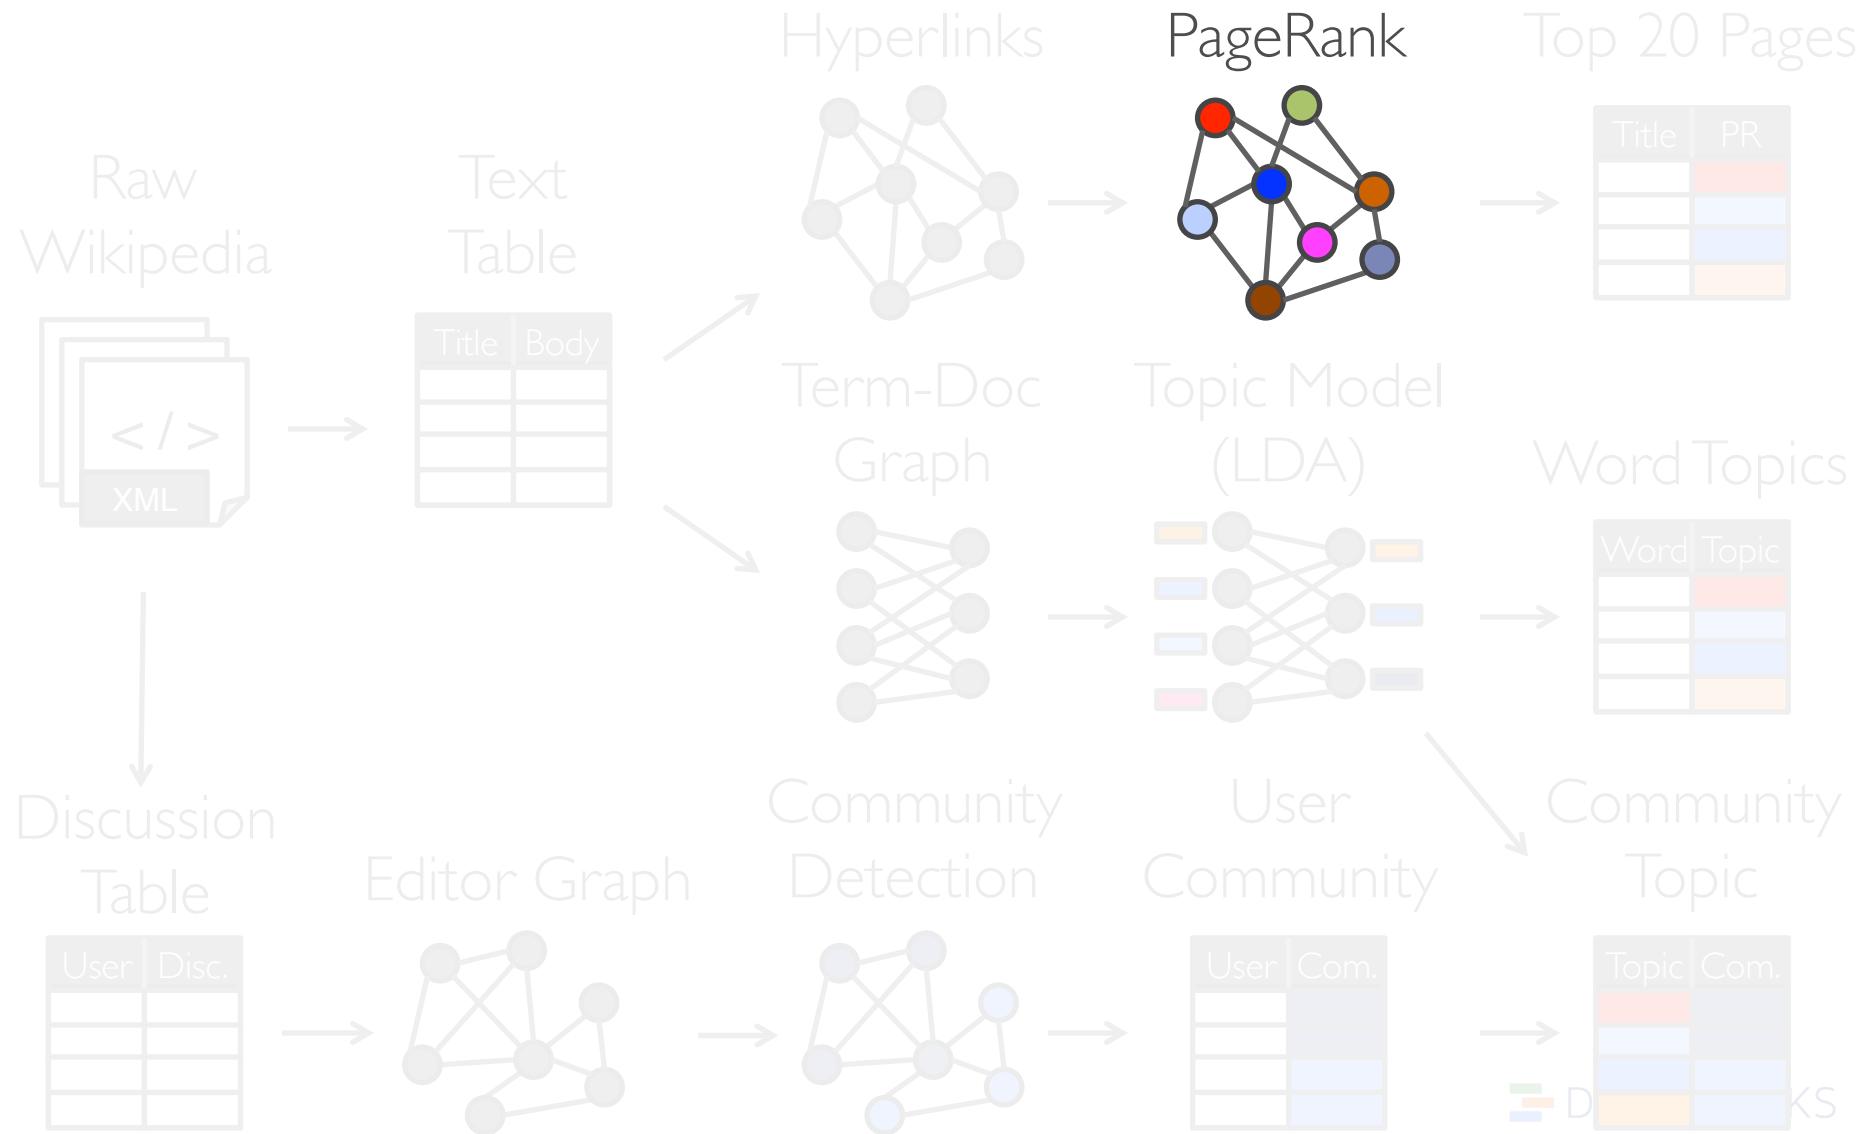
GRAPHX: UNIFIED GRAPH ANALYTICS ON SPARK

Joseph Gonzalez, Reynold Xin, [Ankur Dave](#), Daniel Crankshaw,
Michael Franklin, and Ion Stoica

1. Motivation for GraphX
2. GraphX Implementation
3. Future Directions

Motivation for GraphX

Graphs are Central to Analytics



PageRank: Identifying Leaders

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

Rank of
user i

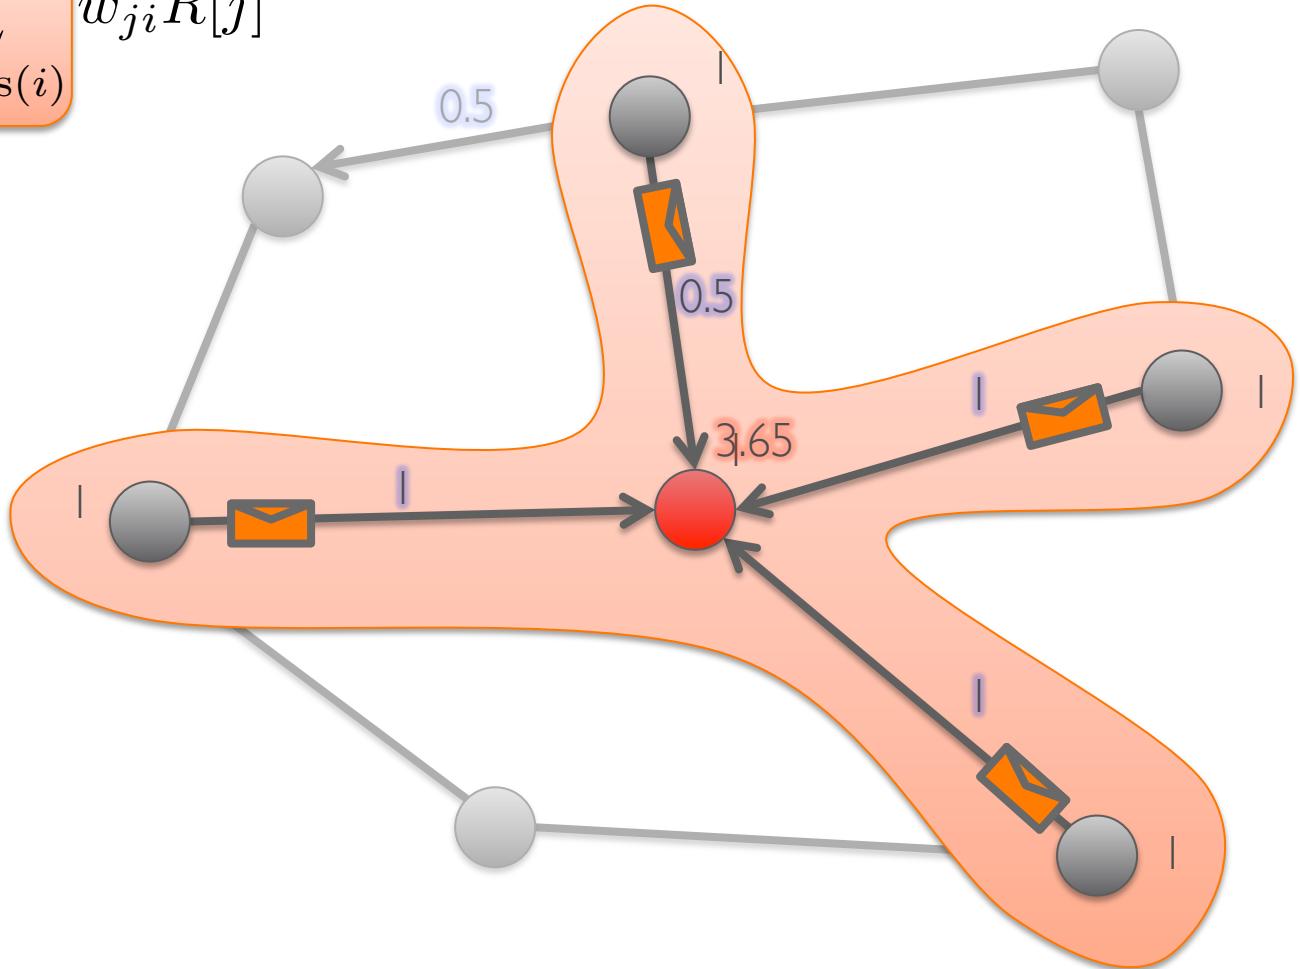
Weighted sum of
neighbors' ranks

Iterate until convergence

Update ranks in parallel

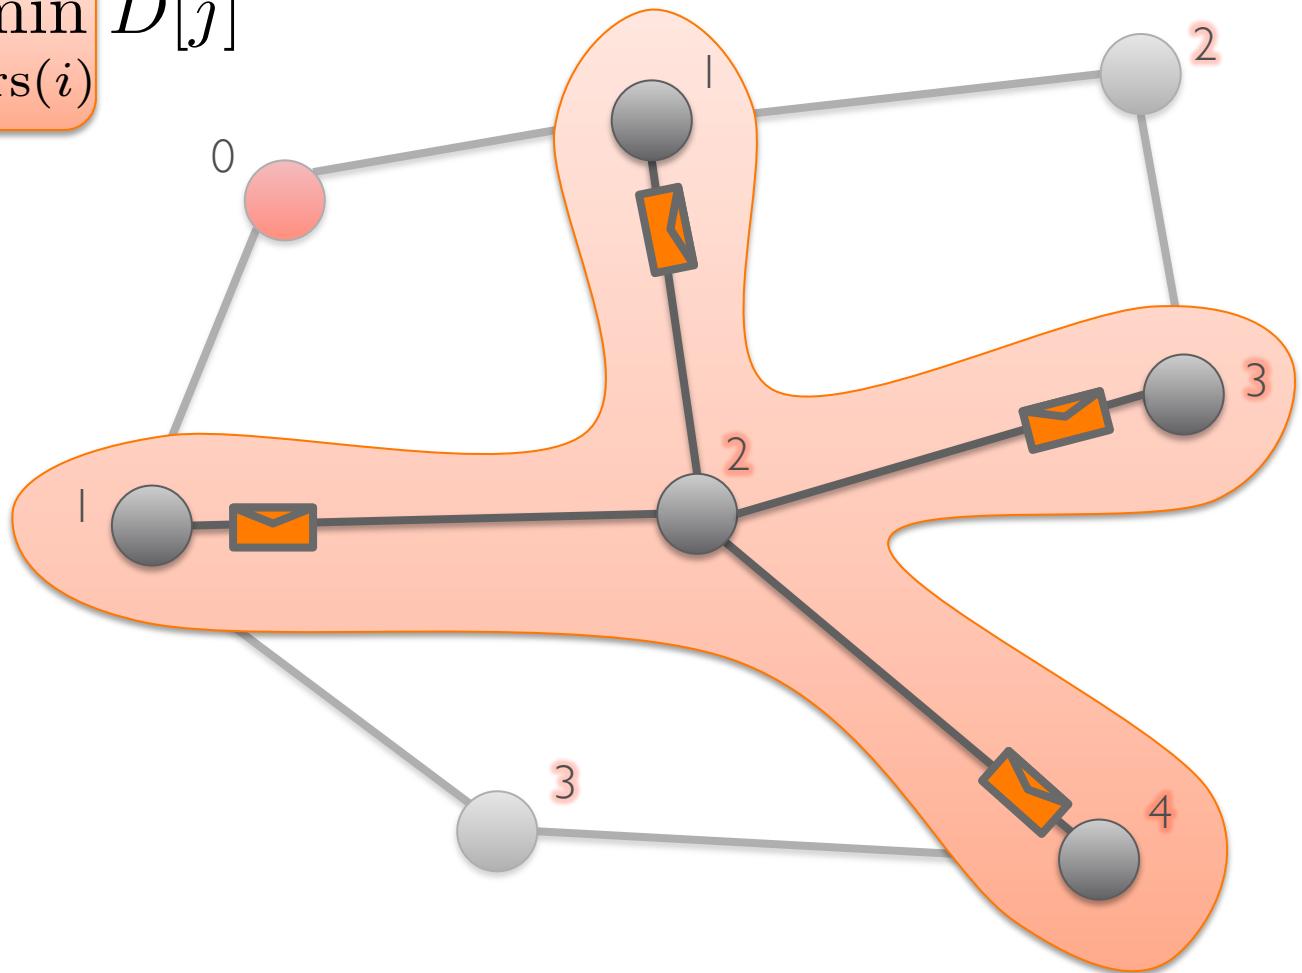
PageRank

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

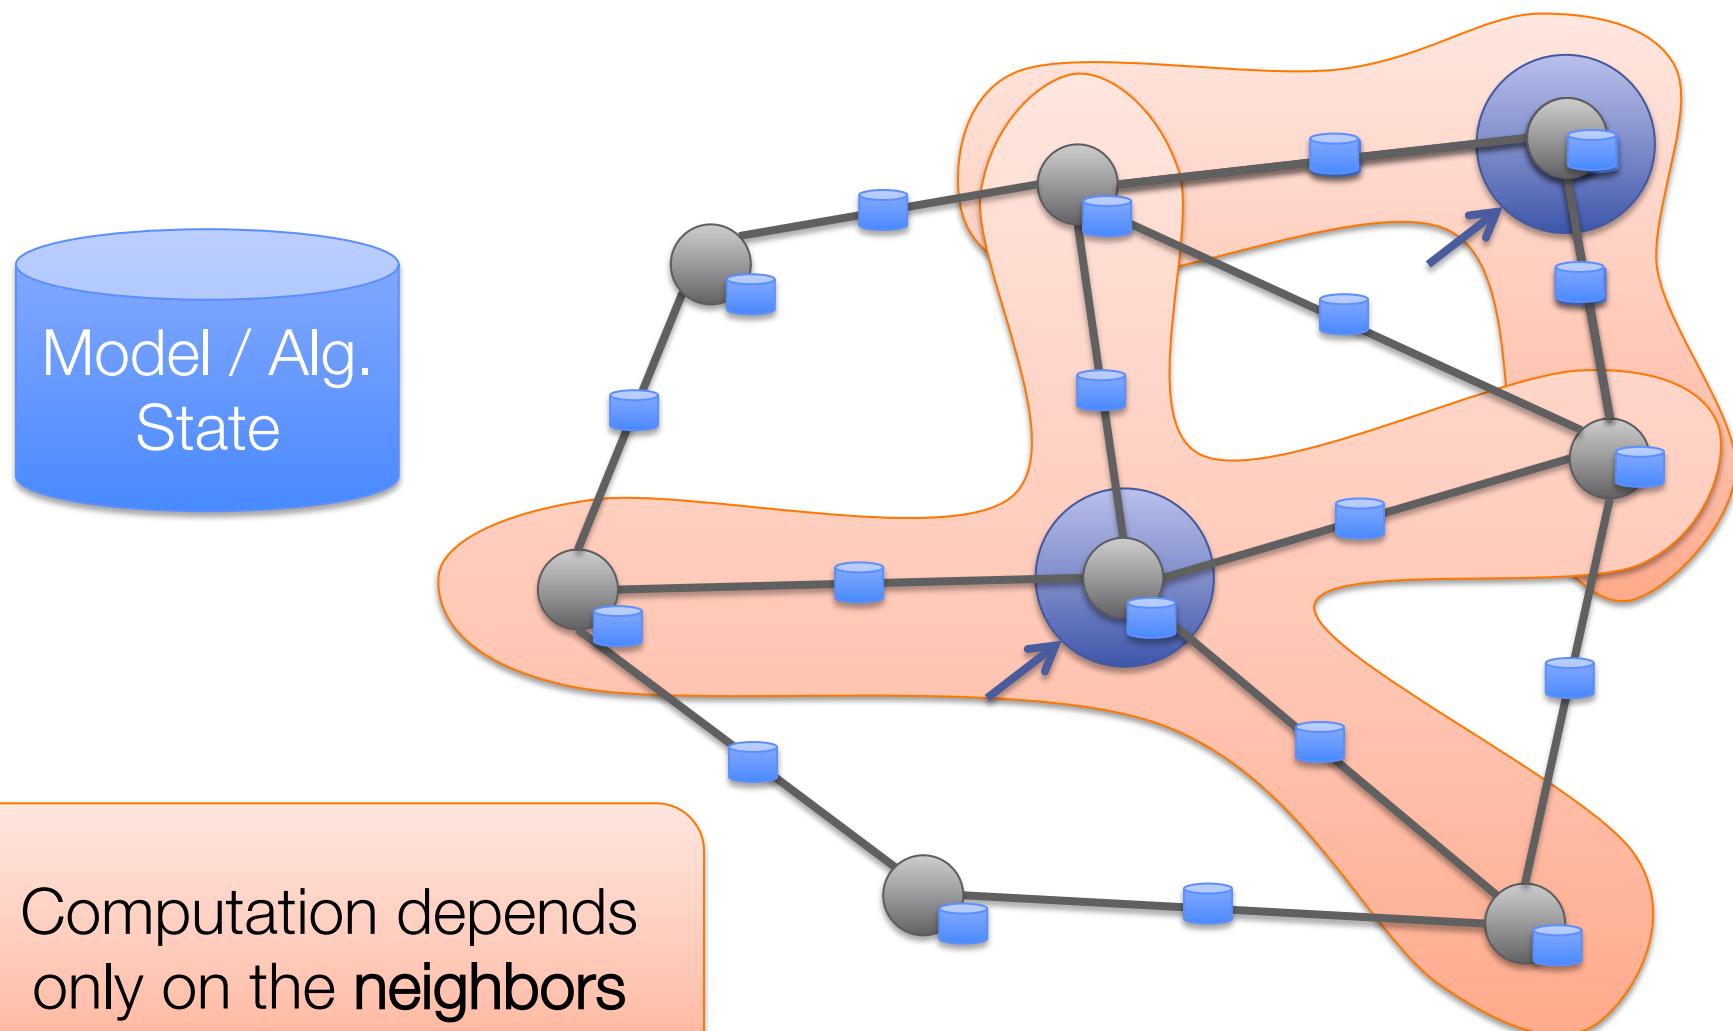


Single-Source Shortest Path

$$D[i] = 1 + \arg \min_{j \in \text{Nbrs}(i)} D[j]$$



The Graph-Parallel Pattern



Many Graph-Parallel Algorithms

Collaborative Filtering

- > Alternating Least Squares
- > Stochastic Gradient Descent
- > Tensor Factorization

Structured Prediction

- > Loopy Belief Propagation
- > Max-Product Linear Programs
- > Gibbs Sampling

Semi-supervised ML

- > Graph SSL
- > CoEM

Community Detection

- > Triangle-Counting
- > K-core Decomposition
- > K-Truss

Graph Analytics

- > PageRank
- > Personalized PageRank
- > Shortest Path
- > Graph Coloring

Classification

- > Neural Networks

Graph-Parallel Systems

Pregel
oo^{gle}



GraphLab

Expose *specialized APIs* to simplify graph programming.

Exploit graph structure to achieve *orders-of-magnitude performance gains* over more general data-parallel systems.

Specialized API: Pregel

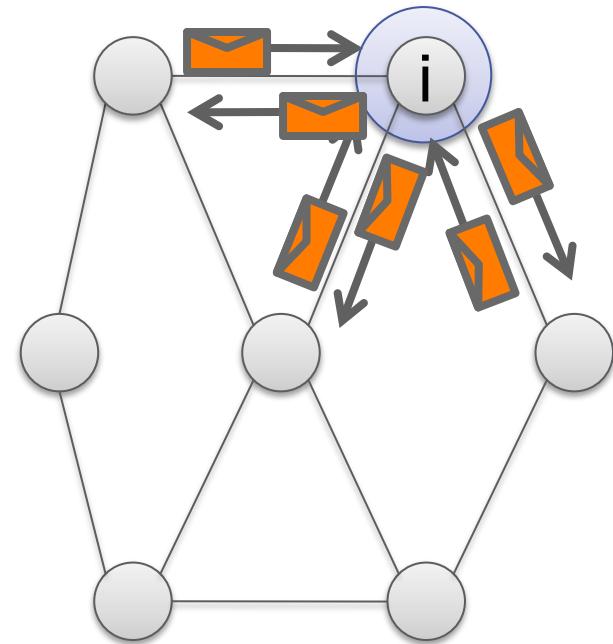
Vertex-Programs interact by sending messages.

```
Pregel_PageRank(i, messages) :
```

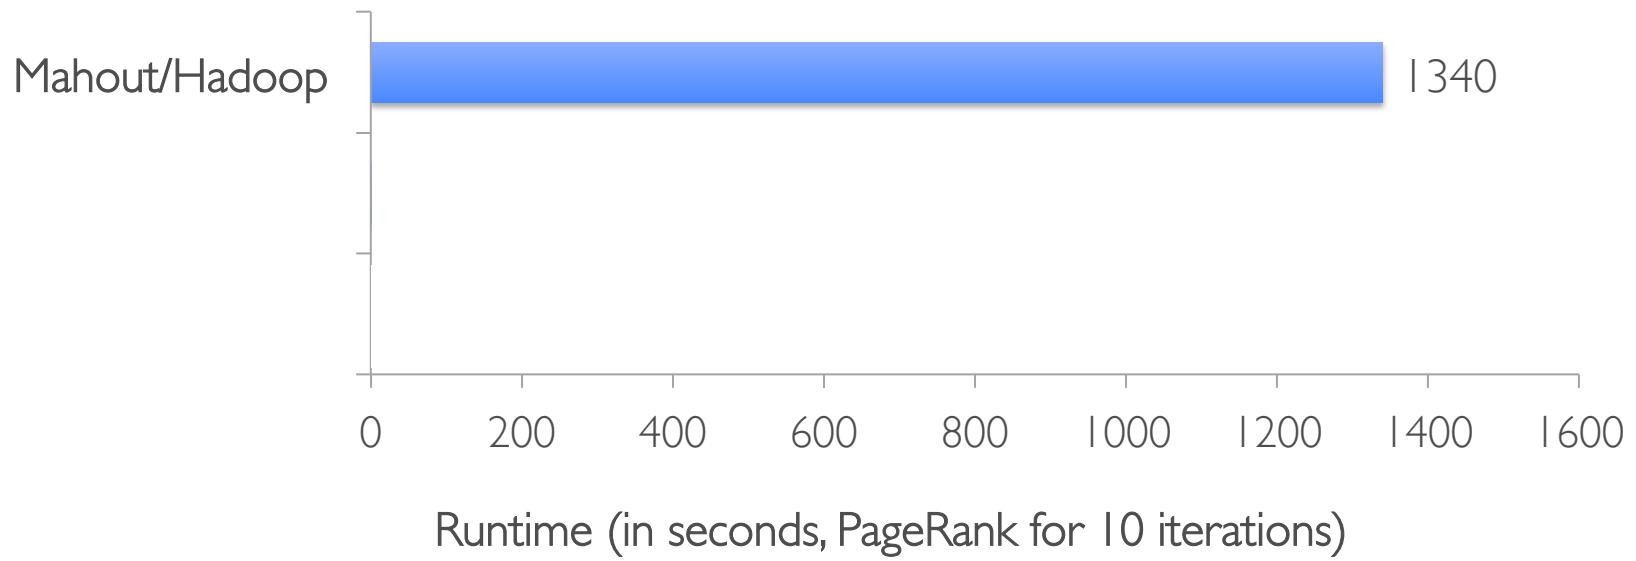
```
    // Receive all the messages  
    total = 0  
    foreach( msg in messages ) :  
        total = total + msg
```

```
    // Update the rank of this vertex  
    R[i] = 0.15 + total
```

```
    // Send new messages to neighbors  
    foreach(j in out_neighbors[i]) :  
        Send msg(R[i]) to vertex j
```

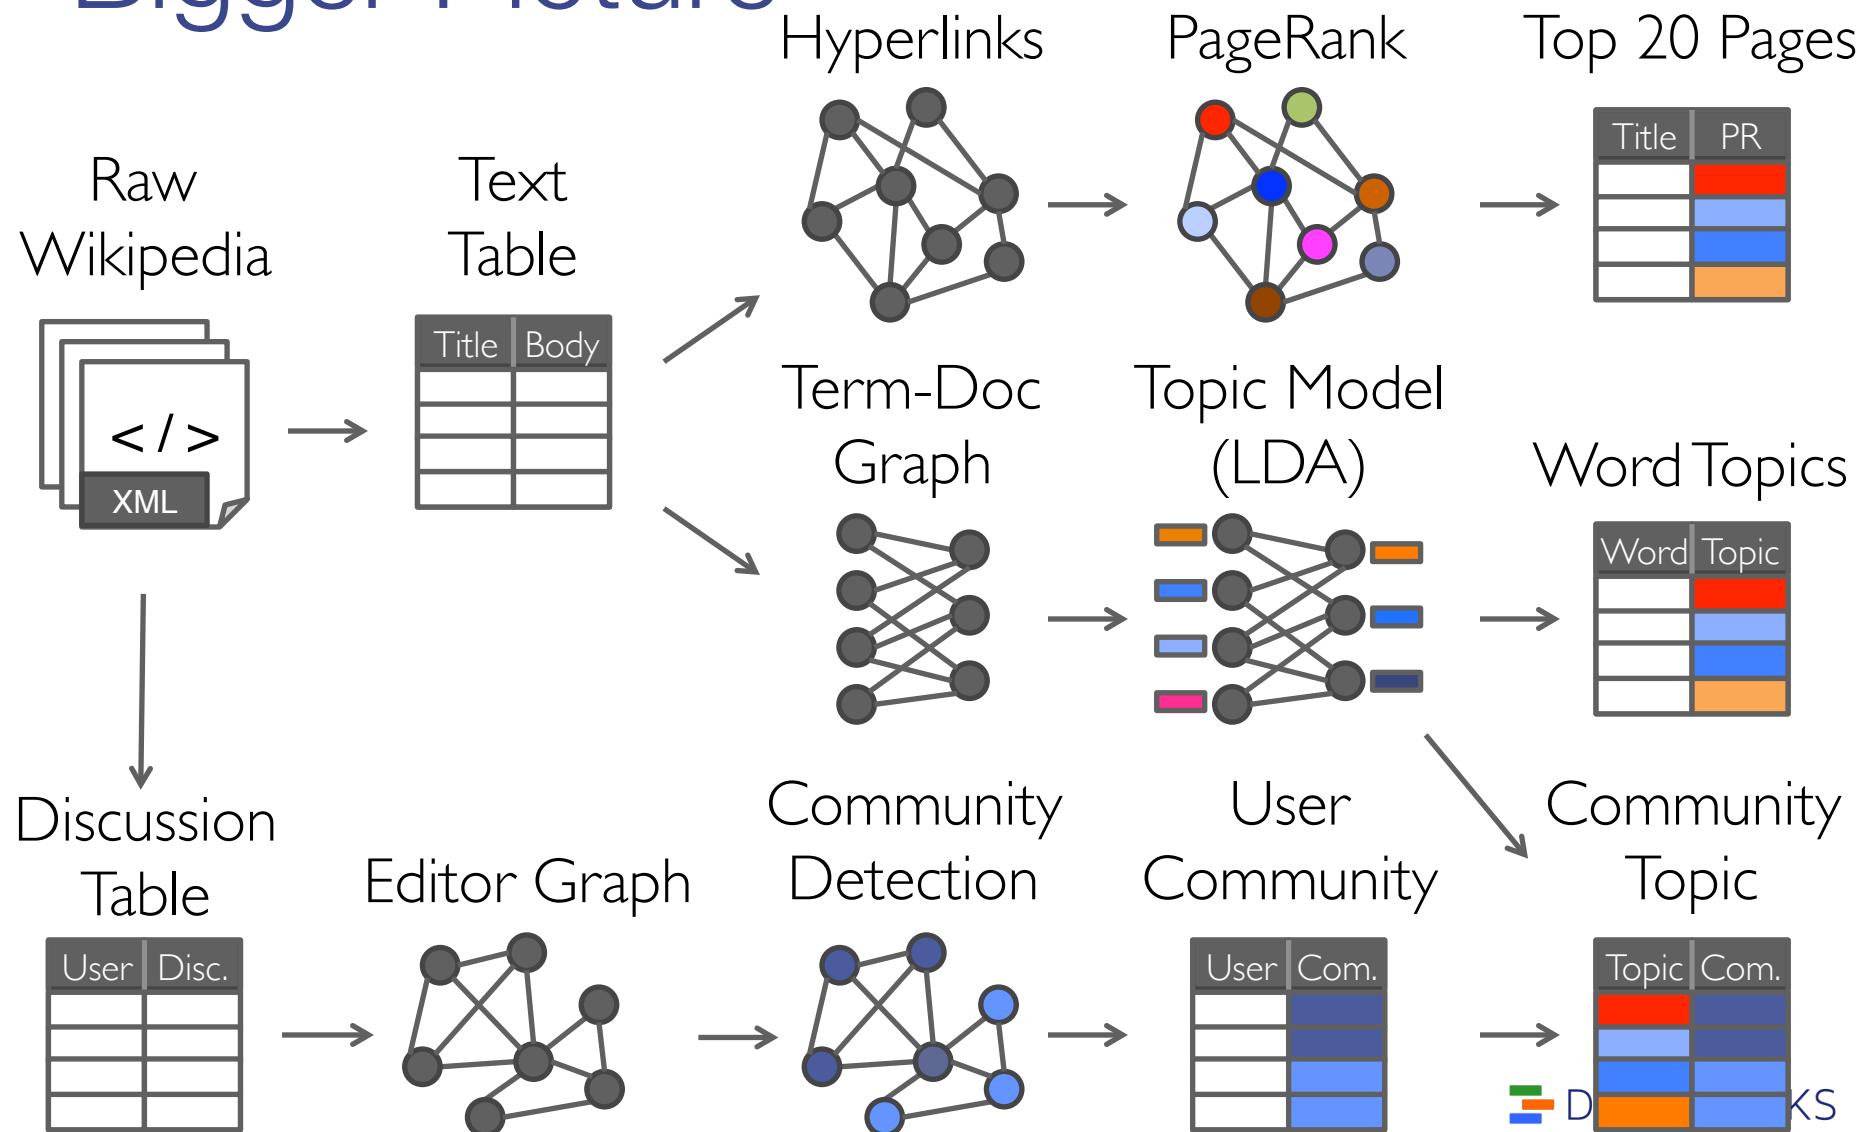


PageRank on LiveJournal Graph (69M edges)



Spark is *4x faster* than Hadoop
GraphLab is *16x faster* than Spark

Specialized Systems Miss the Bigger Picture



Separate Systems to Support Each View

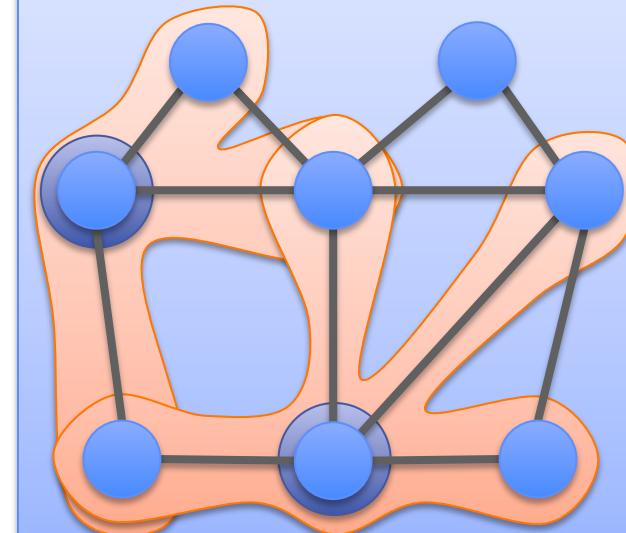
Data-Parallel



Graph-Parallel



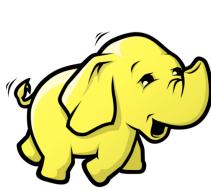
Dependency Graph



*Having separate systems
for each view is
difficult to use and inefficient*

Difficult to Program and Use

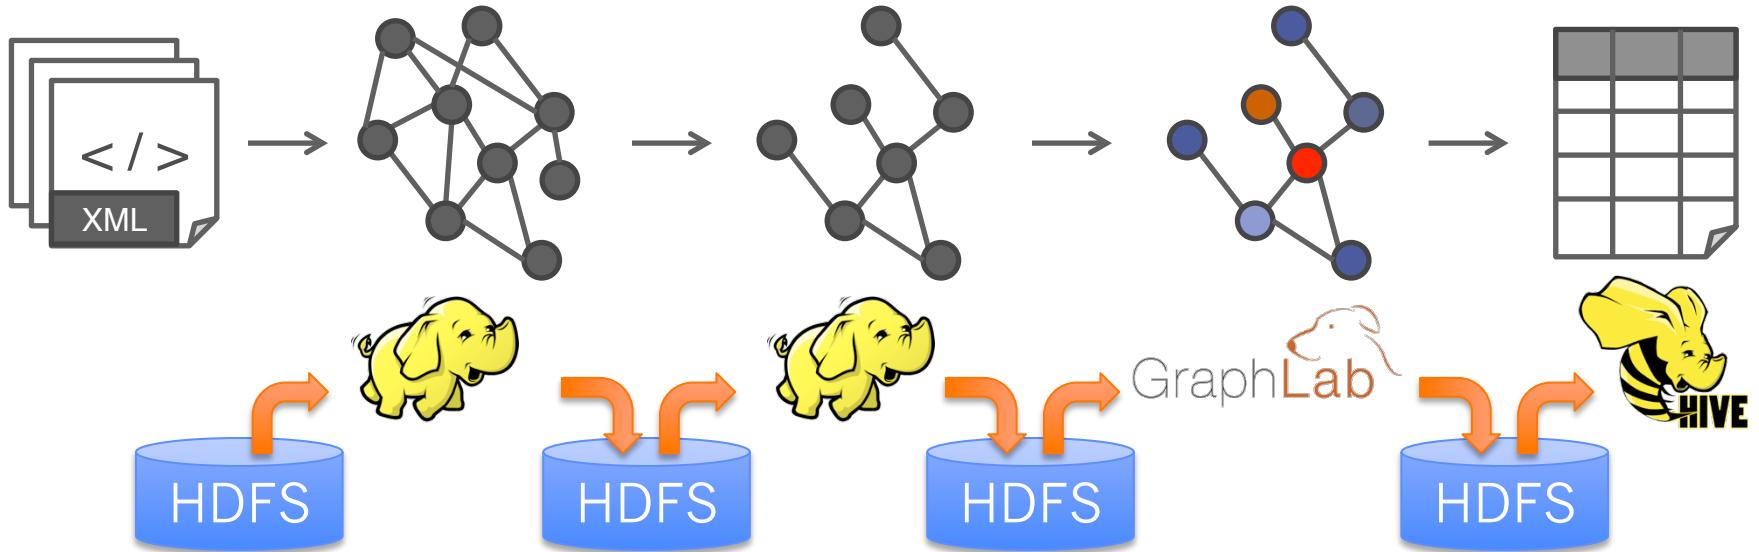
Users must *Learn*, *Deploy*, and *Manage*
multiple systems



Leads to brittle and often
complex interfaces

Inefficient

Extensive **data movement** and **duplication** across
the network and file system

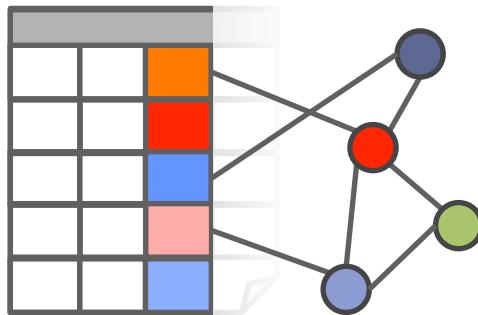


Limited reuse internal data-structures
across stages

Solution: The GraphX Unified Approach

New API

*Blurs the distinction between
Tables and Graphs*



New Library

*Embeds Graph-Parallel
model in Spark*



Enabling users to **easily** and **efficiently**
express the entire graph analytics pipeline

Tables and Graphs are composable views of the same physical data

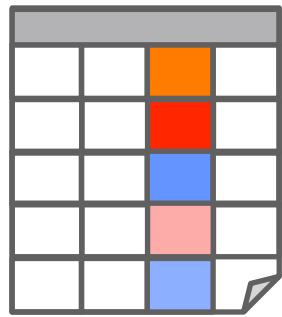
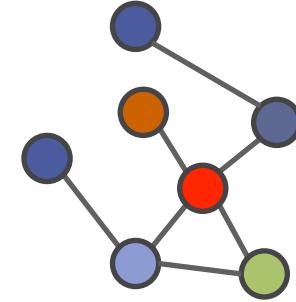
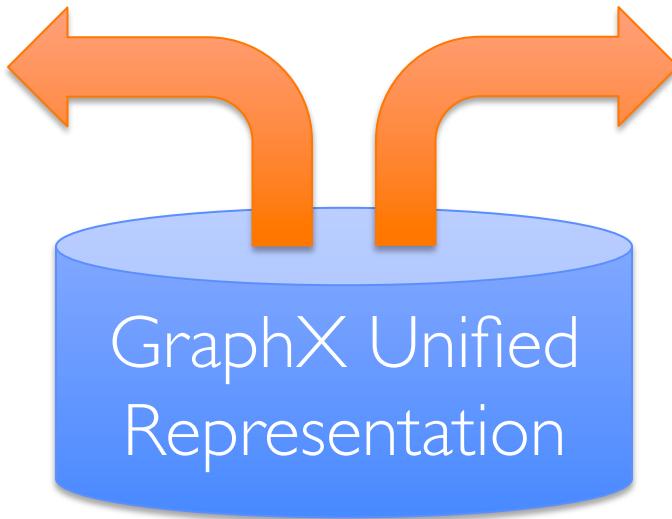


Table View

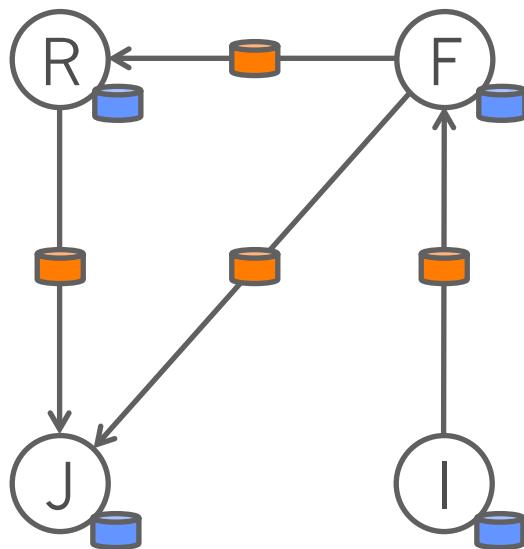


Graph View

Each view has its own operators that
exploit the semantics of the view
to achieve efficient execution

View a Graph as a Table

Property Graph



Vertex Table

Id	Attribute (V)
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk)
Istoica	(Prof., Berk)

Edge Table

SrcId	DstId	Attribute (E)
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI

Table Operators

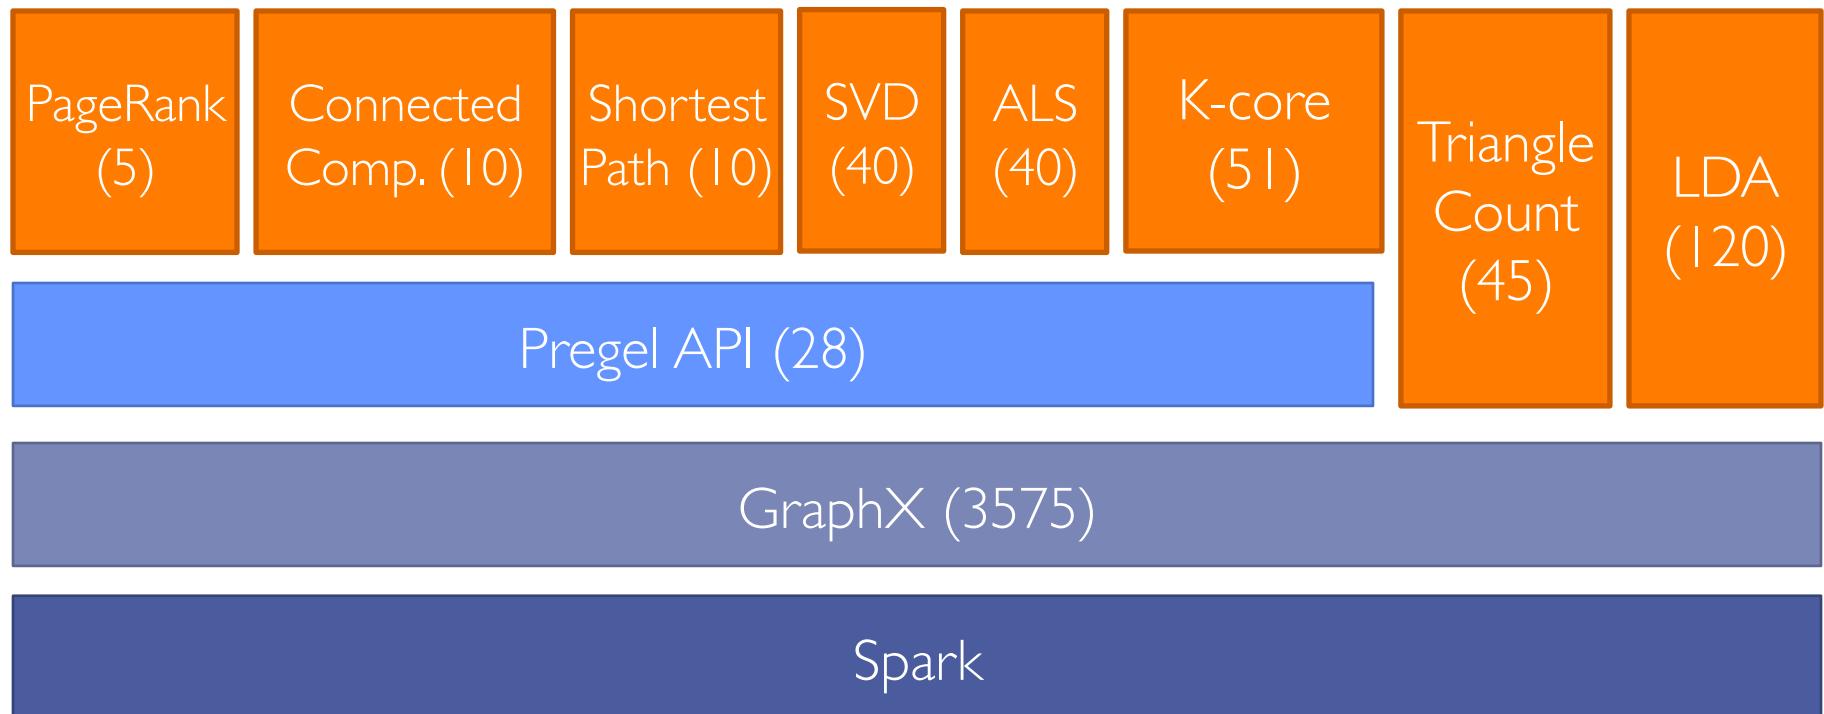
Table (RDD) operators are inherited from Spark:

map	reduce	sample
filter	count	take
groupBy	fold	first
sort	reduceByKey	partitionBy
union	groupByKey	mapwith
join	cogroup	pipe
leftOuterJoin	cross	save
rightOuterJoin	zip	...

Graph Operators

```
class Graph [ V, E ] {  
    def Graph(vertices: Table[ (Id, V) ],  
              edges: Table[ (Id, Id, E) ])  
        // Table views -----  
        def vertices: Table[ (Id, V) ]  
        def edges: Table[ (Id, Id, E) ]  
        def triplets: Table [ ((Id, V), (Id, V), E) ]  
        // Computation -----  
        def mrTriplets(mapF: (Edge[V, E]) => List[(Id, T)],  
                      reduceF: (T, T) => T): Graph[T, E]  
  
        // Convenience functions -----  
        def mapV(m: (Id, V) => T ): Graph[T, E]  
        def mapE(m: Edge[V, E] => T ): Graph[V, T]  
        def joinV(tbl: Table [(Id, T)]): Graph[(V, T), E ]  
        def joinE(tbl: Table [(Id, Id, T)]): Graph[V, (E, T)]  
        def reverse: Graph[V, E]  
        def subgraph(pv: (Id, V) => Boolean,  
                    pe: Edge[V, E] => Boolean): Graph[V, E]
```

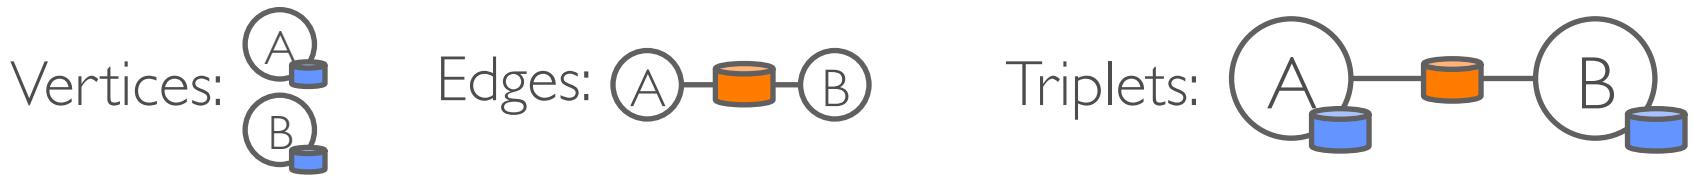
The GraphX Stack (Lines of Code)



Triplets Join Vertices and Edges

The *triplets* operator joins vertices and edges:

```
SELECT src.Id, dst.Id, src.attr, e.attr, dst.attr  
FROM edges AS e JOIN vertices AS src, vertices AS dst  
ON e.srcId = src.Id AND e.dstId = dst.Id
```



The *mrTriples* operator sums adjacent triplets.

```
SELECT t.dstId, reduceUDF( mapUDF(t) ) AS sum  
FROM triplets AS t GROUPBY t.dstId
```

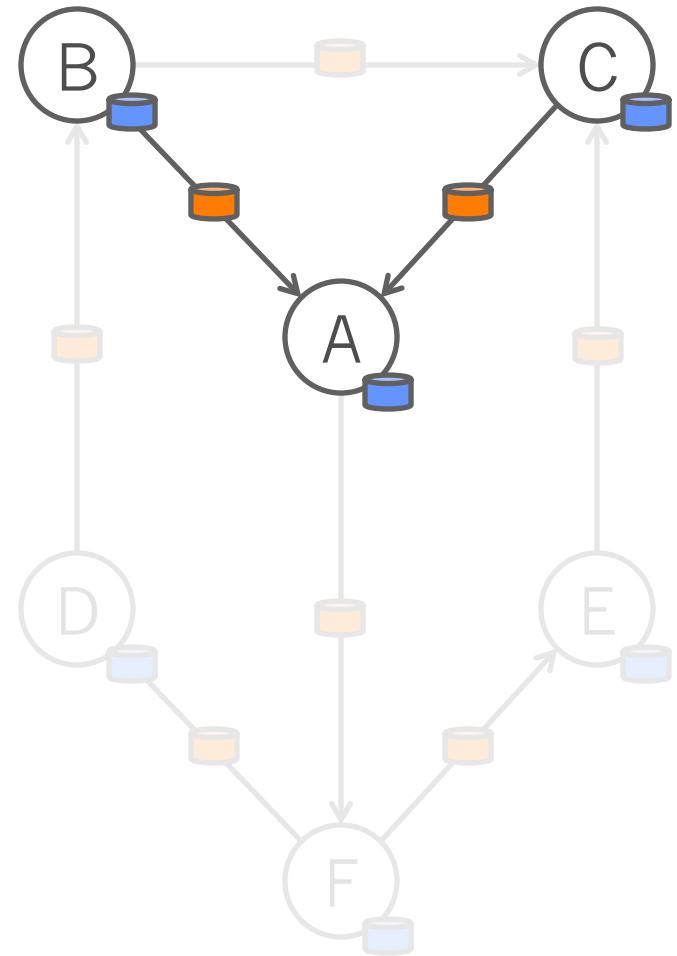
Map Reduce Triplets

Map-Reduce for each vertex

mapF($(A \leftarrow B)$) $\rightarrow A_1$

mapF($(A \leftarrow C)$) $\rightarrow A_2$

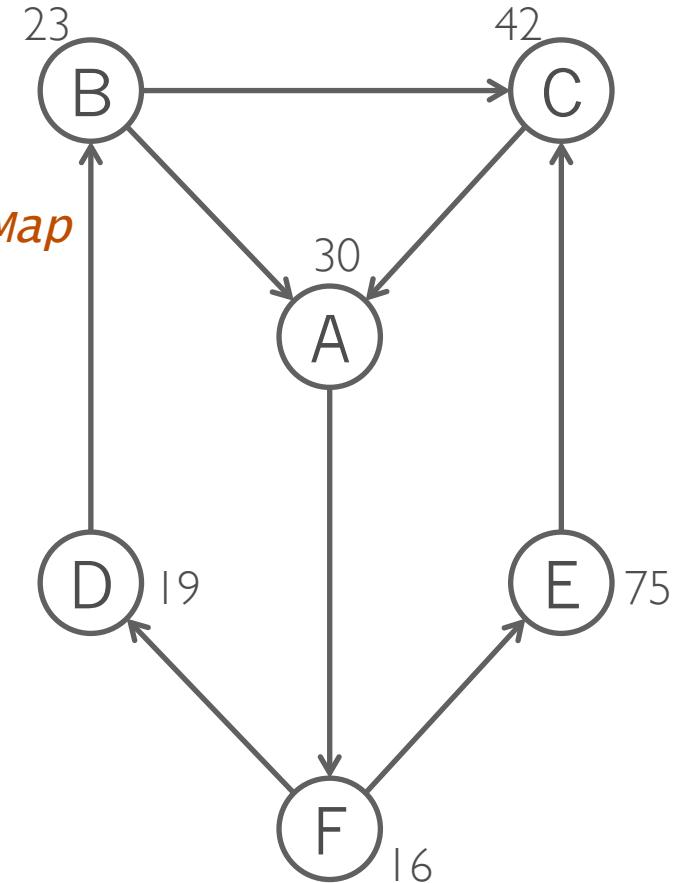
reduceF((A_1, A_2)) $\rightarrow A$



Example: Oldest Follower

What is the age of the oldest follower for each user?

```
val oldestFollowerAge = graph  
  .mrTriplets(  
    e=> (e.dst.id, e.src.age), //Map  
    (a,b)=> max(a, b) //Reduce  
  )  
  .vertices
```



Graphs are
first-class
objects

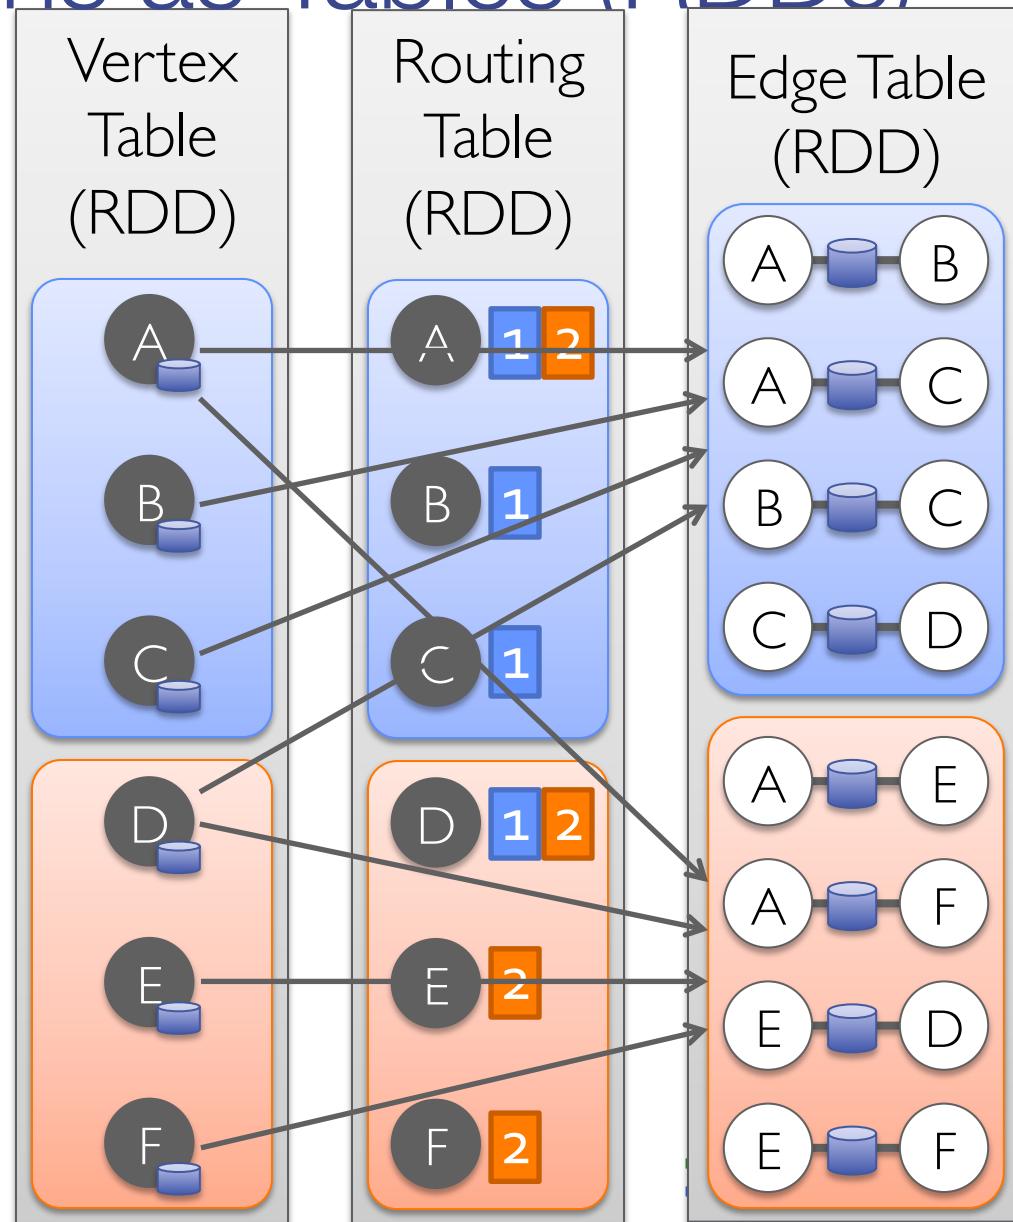
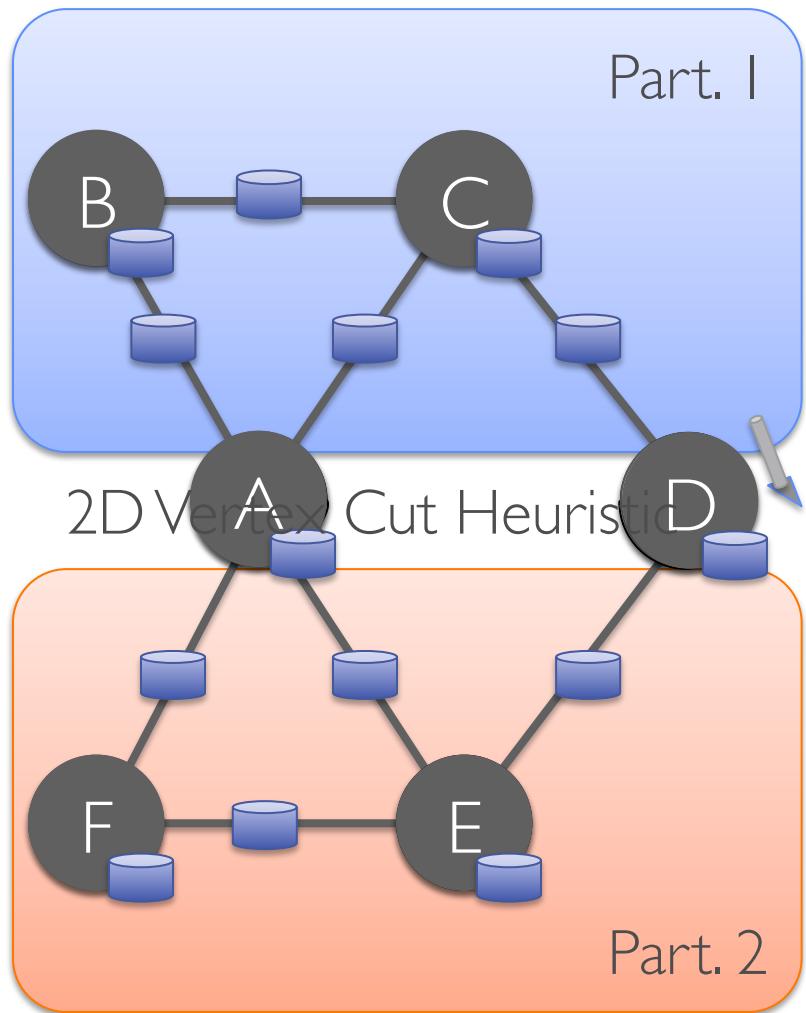
Rank in GraphX

```
// Load and initialize the graph
val graph = GraphLoader.edgeListFile("hdfs://web.txt")
val prGraph = graph.joinVertices(graph.outDegrees)
// Implement and Run PageRank
val pageRank =
  prGraph.pregel(initialMessage = 0.0, iter = 10)(
    (oldV, msgSum) => 0.15 + 0.85 * msgSum,
    triplet => triplet.src.pr / triplet.src.deg,
    (msgA, msgB) => msgA + msgB)
// Get the top 20 pages
pageRank.vertices.top(20)(Ordering.by(_.pr)).foreach(println)
```

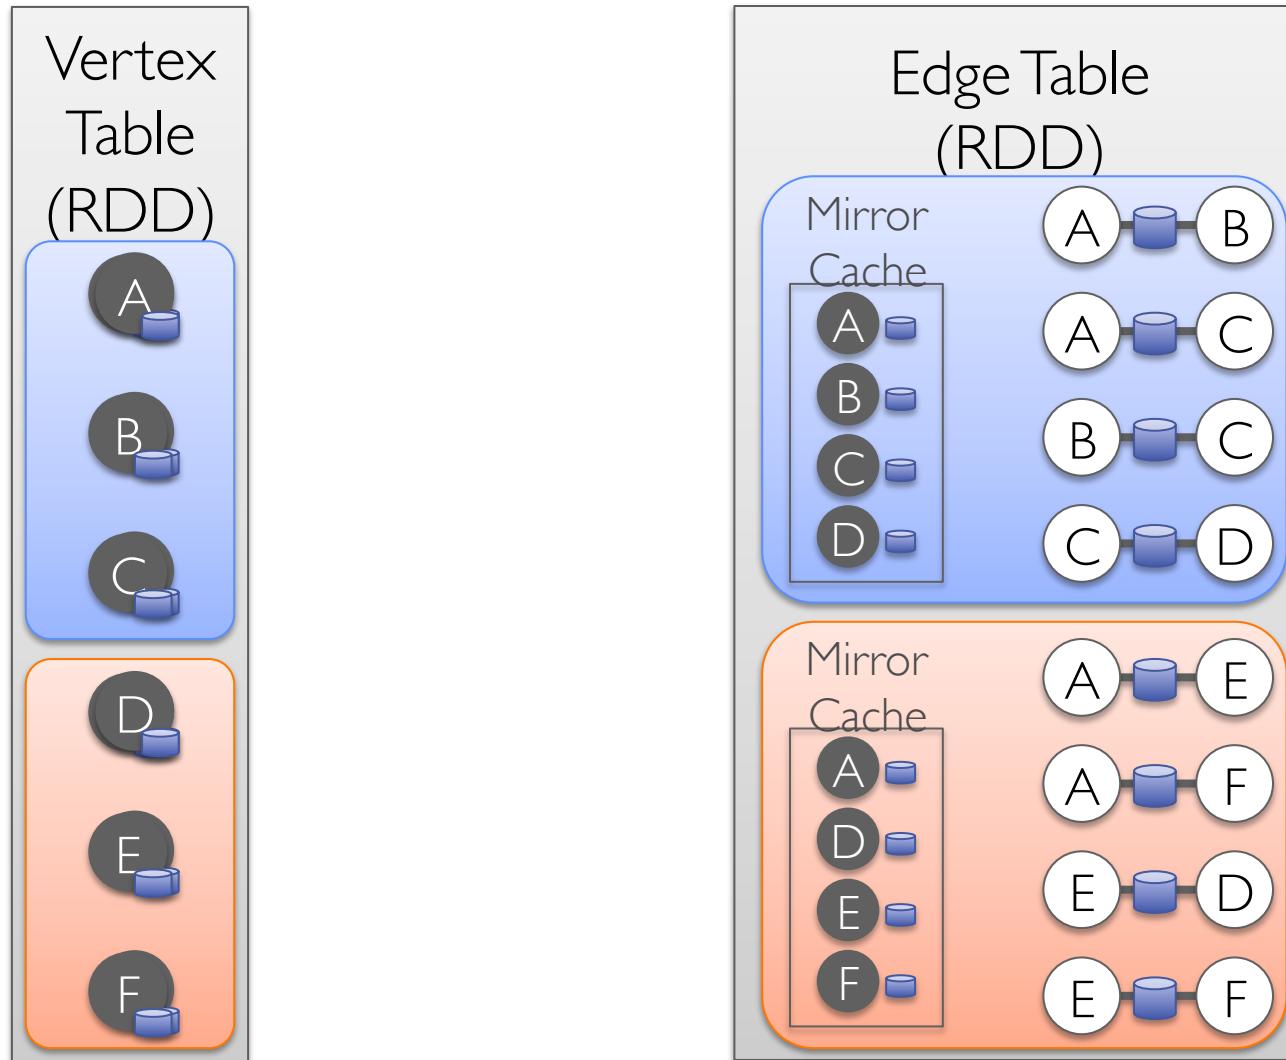
Implementation

Distributed Graphs as Tables (RDDs)

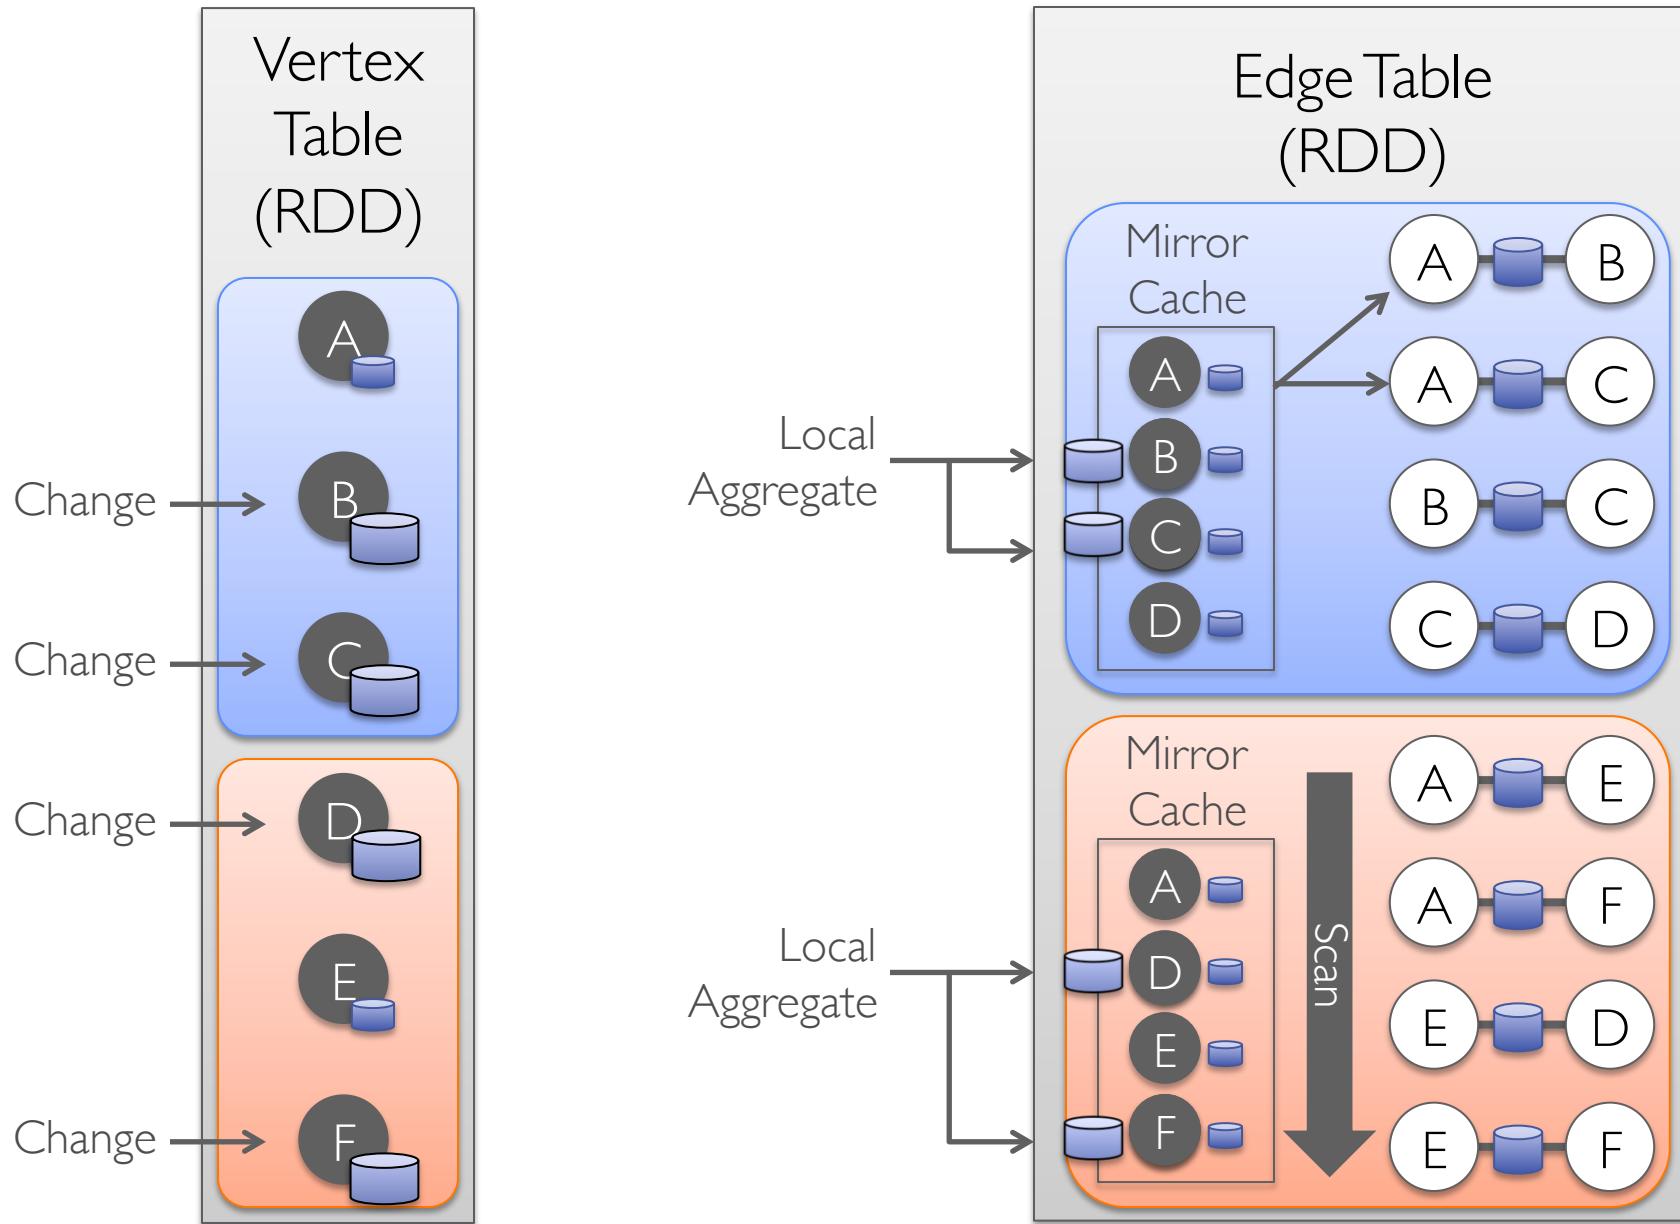
Property Graph



Caching for Iterative mrTriplets



mrTriplets Execution

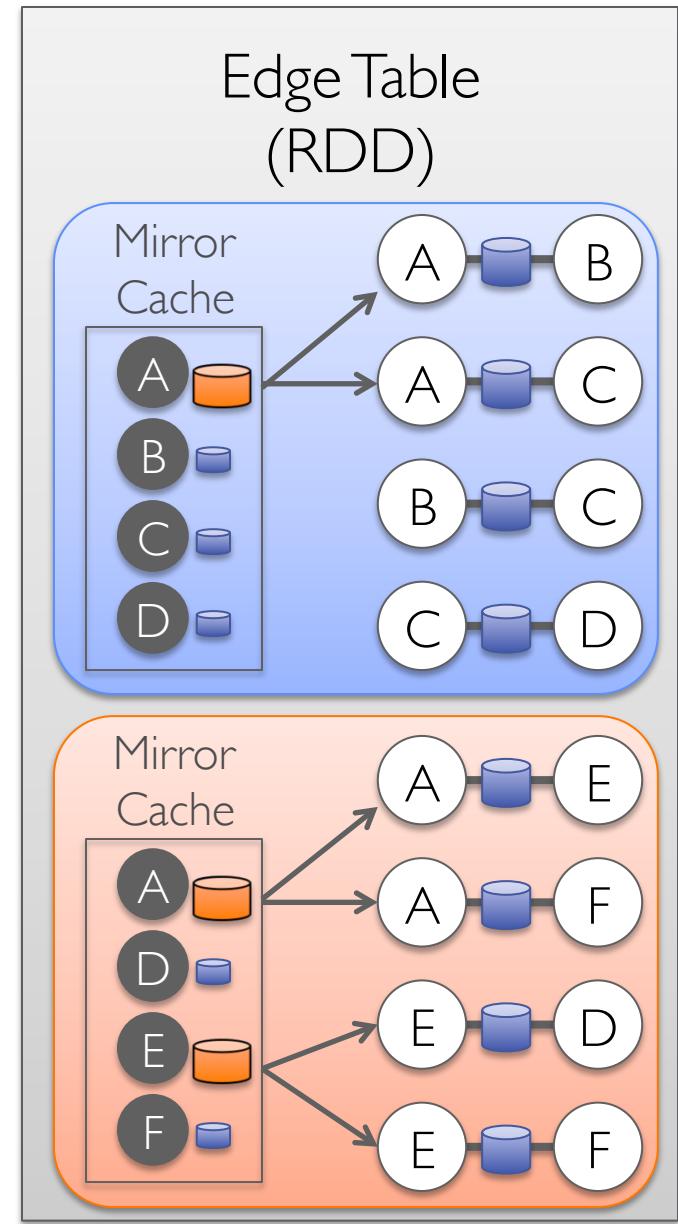
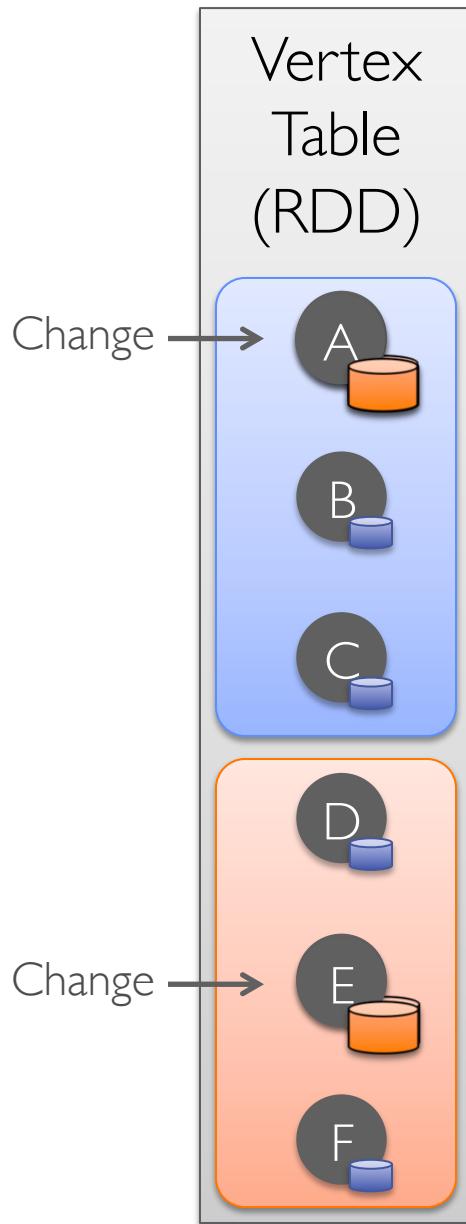


BRICKS

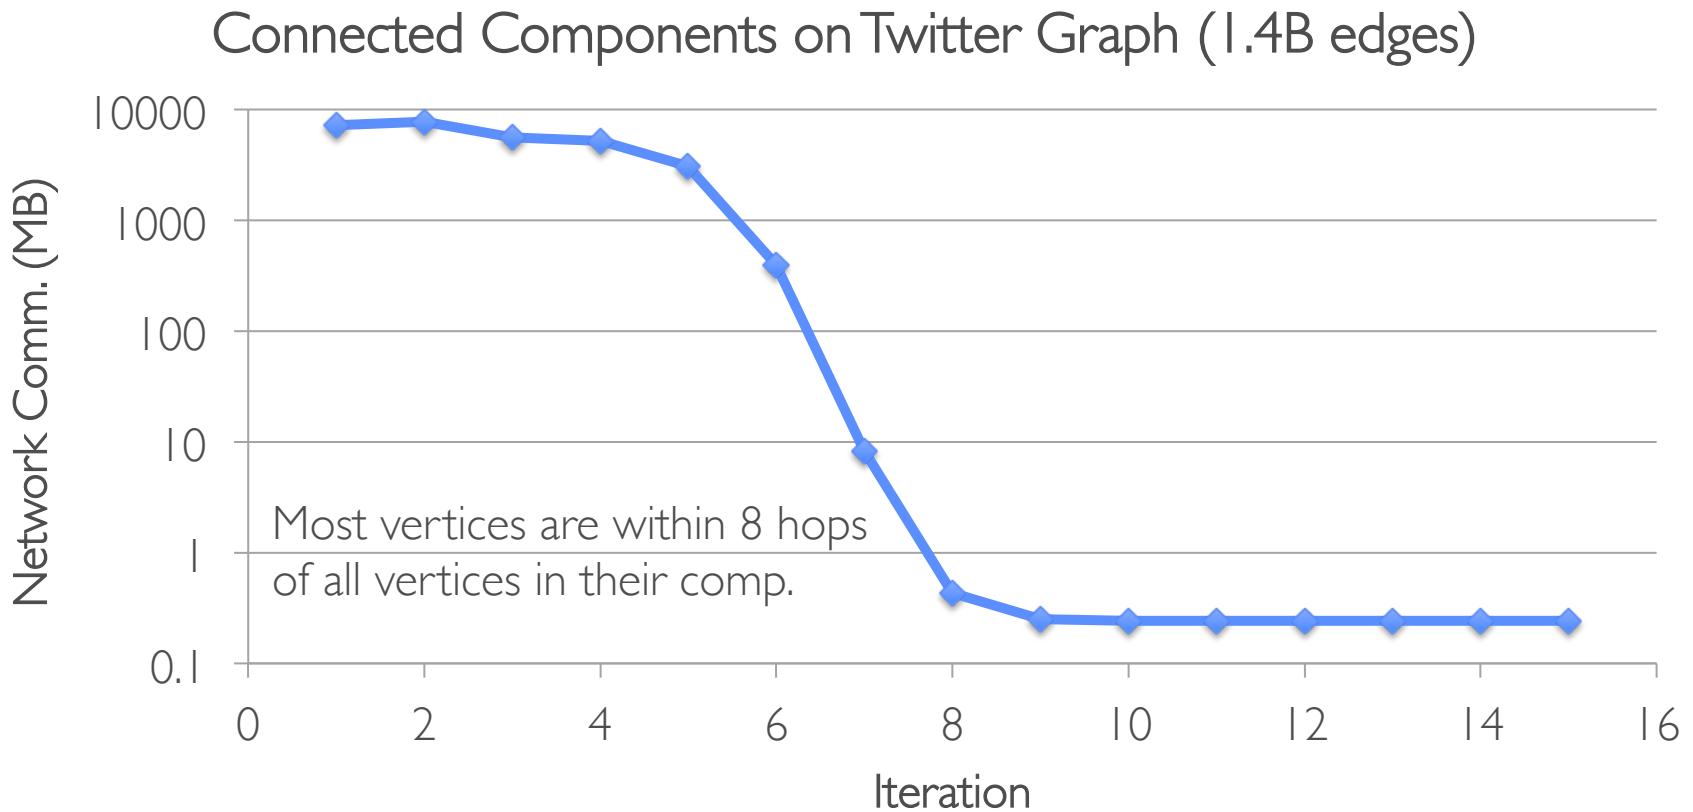
Optimizations

1. Incremental Updates to Mirror Caches
2. Join Elimination
3. Index Scanning for Active Sets
4. Local Vertex and Edge Indices
5. Index and Routing Table Reuse

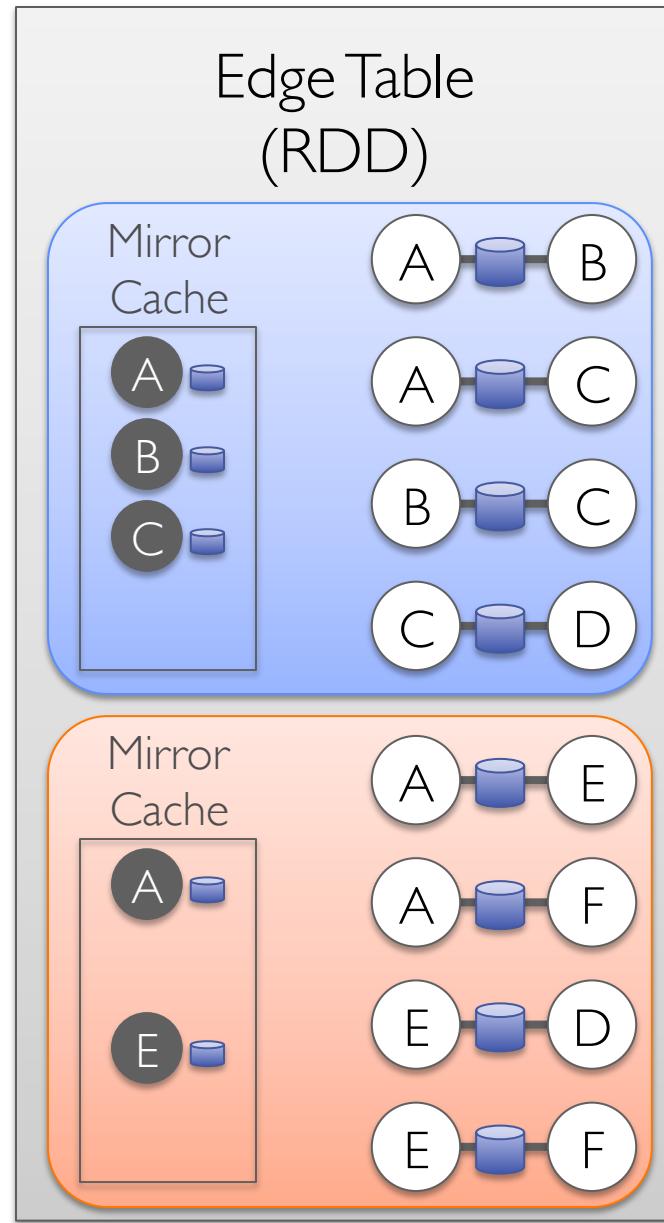
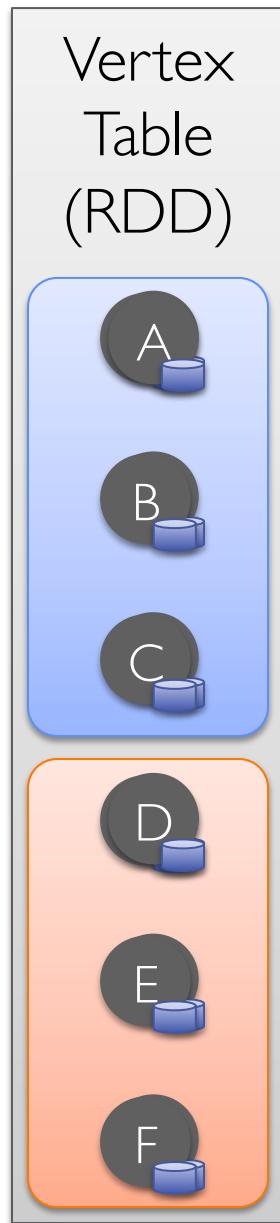
I. Incremental Updates for Mirror Caches



Reduction in Communication Due to Cached Updates



2. Join Elimination

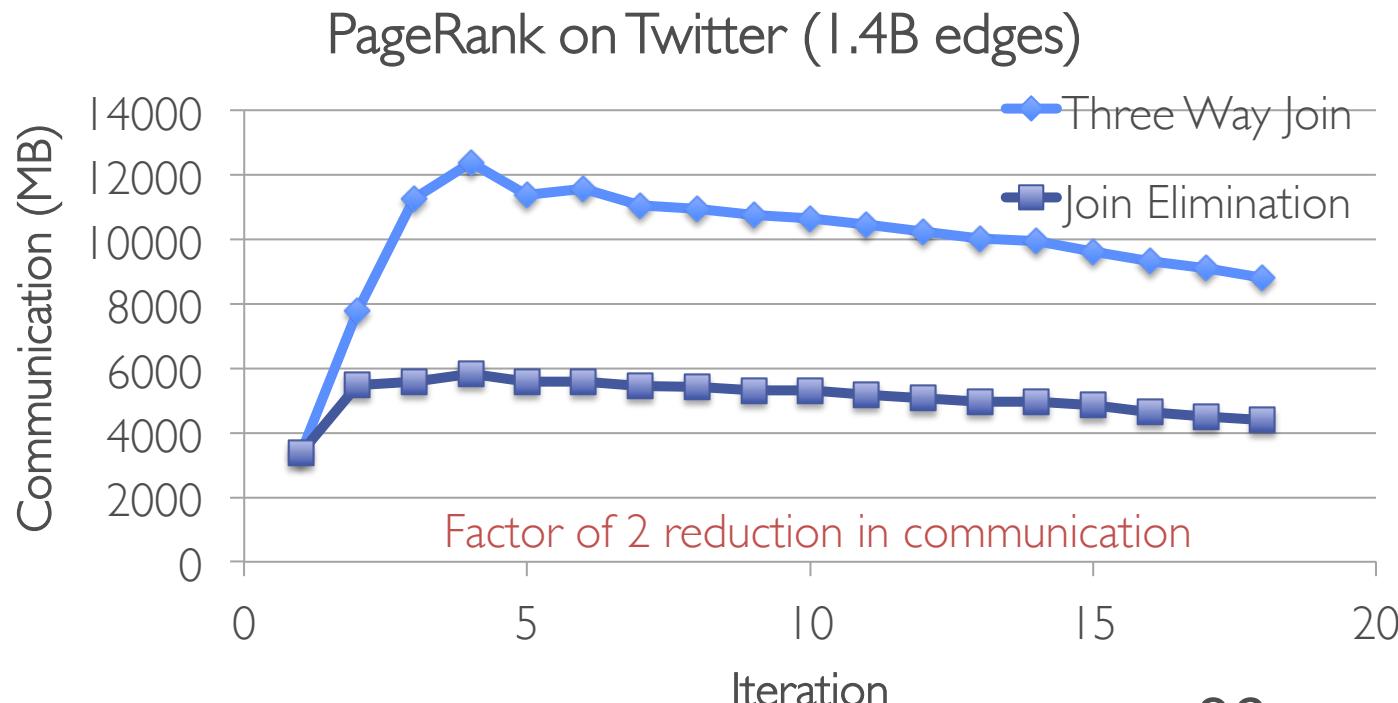


BRICKS

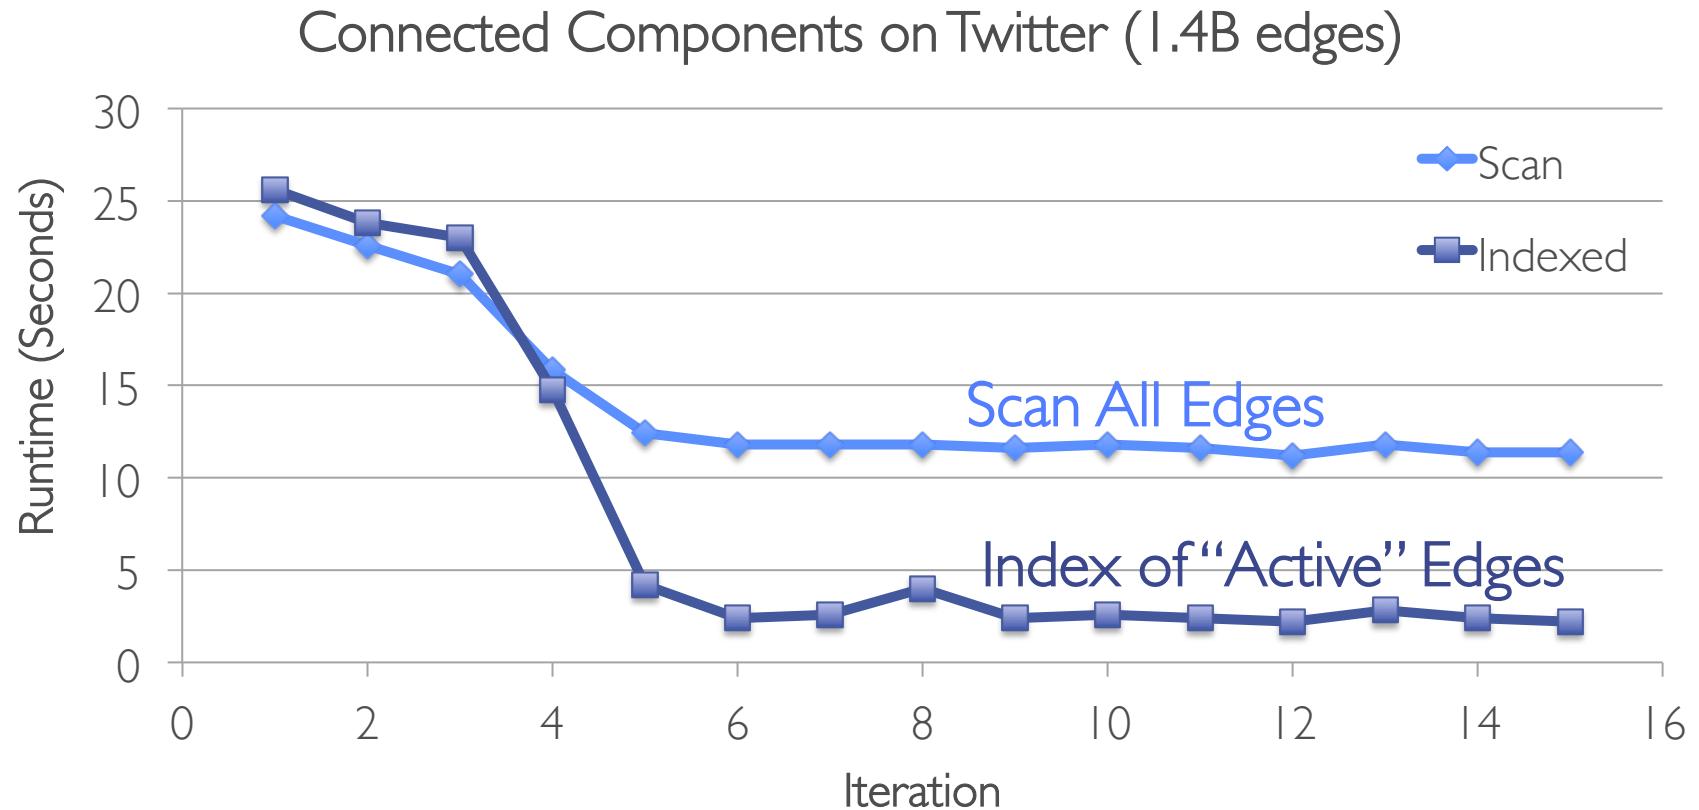
Reduction in Communication Due to Join Elimination

UDF byte code inspection for join elimination:

- > Identify and bypass joins for unused triplets fields
- > *Example:* PageRank only accesses source attribute



3. Index Scanning for Active Sets



Additional Query Optimizations

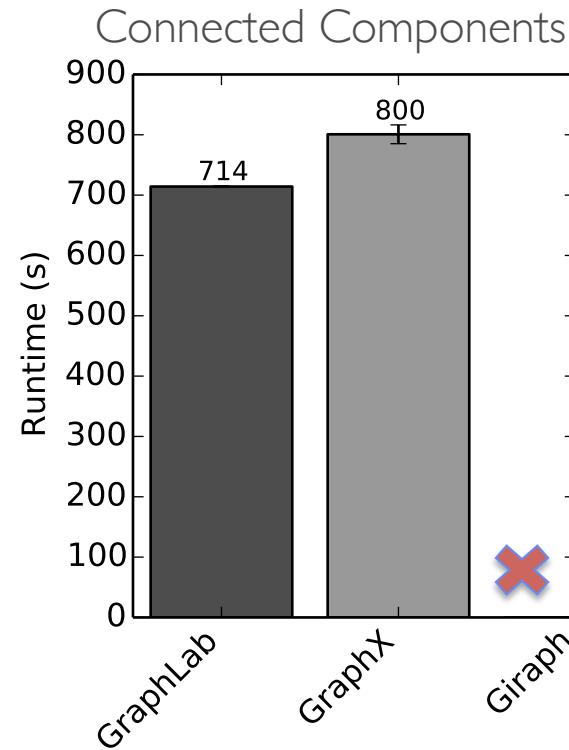
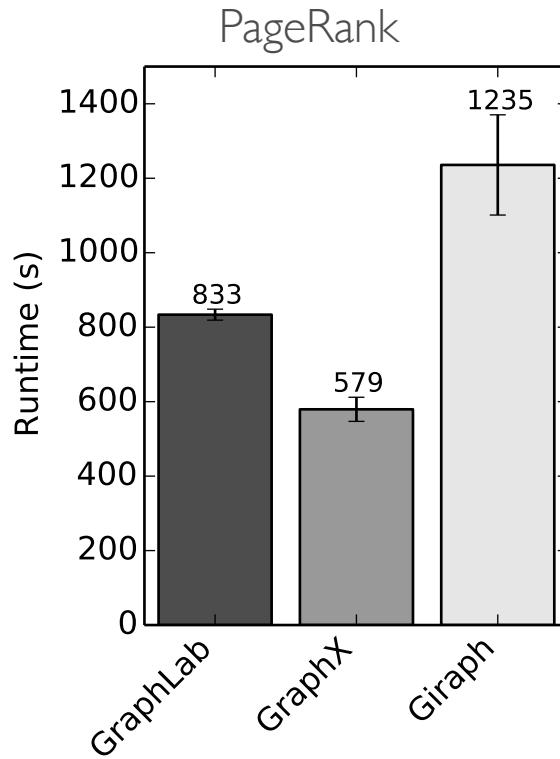
Indexing and Bitmaps:

- > To accelerate joins across graphs
- > To efficiently construct sub-graphs

Substantial Index and Data Reuse:

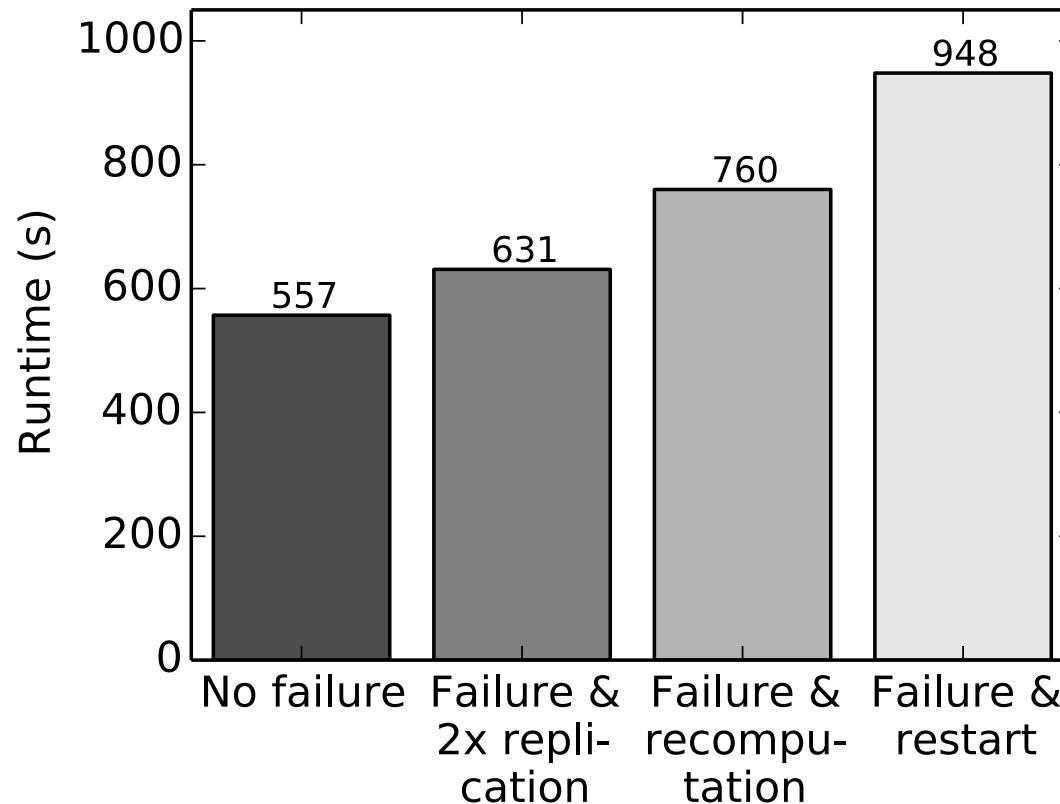
- > *Split Hash Table*: reuse key sets across hash tables
- > Reuse routing tables across graphs and sub-graphs
- > Reuse edge adjacency information and indices

Multi-System Comparison



uk-2007-05 web graph (3.7B edges)

Fault Tolerance



GraphX Training

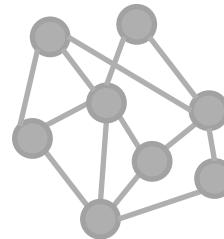
Raw
Wikipedia



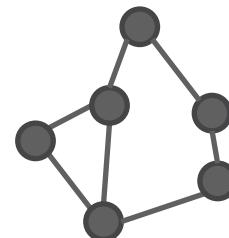
Text
Table

Title	Body

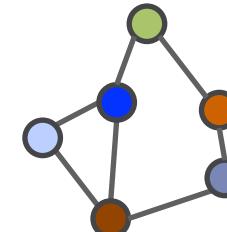
Hyperlinks



Berkeley
subgraph



PageRank



Top 20 Pages

Title	PR
	Red
	Blue
	Blue
	Orange

Demo

Future Directions

- Computation on time-varying graphs
- Graph serving
- Operating on compressed graphs

Thanks!

<http://spark.apache.org/graphx>

{jegonzal, rxin, ankurd, crankshaw}@eecs.berkeley.edu